

Networking lecture 2

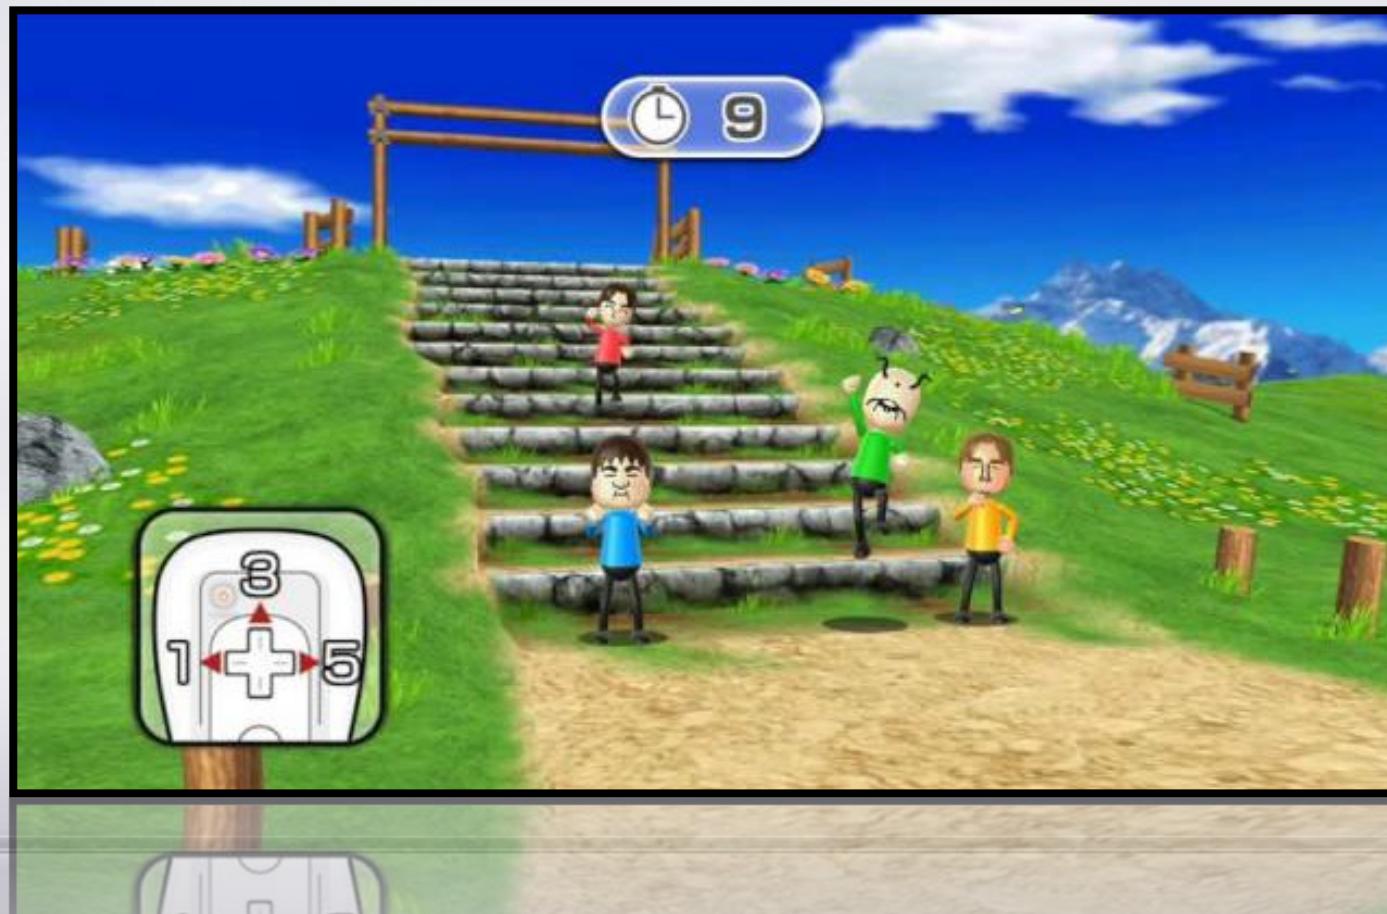
- UDP vs TCP
- (TCP) networking challenges
- Assignment 2

Recap

"Uhm... networking right?"

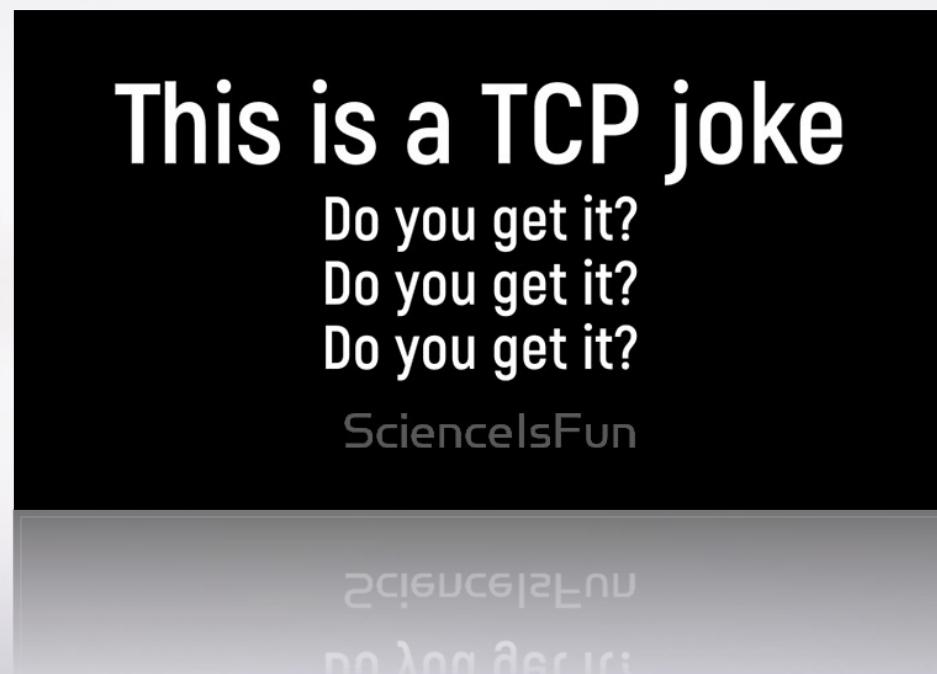
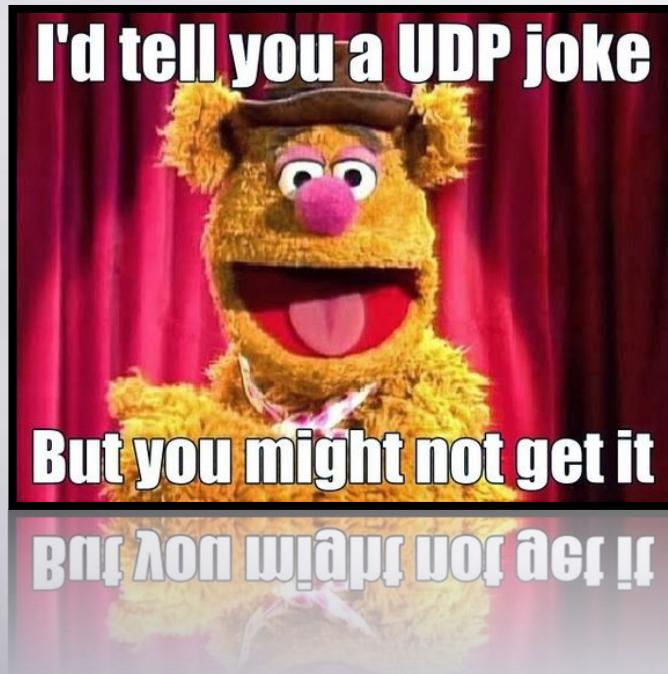
Recap 1 – Why are we here again?

"To build a (simple) network game on top of low level technologies"



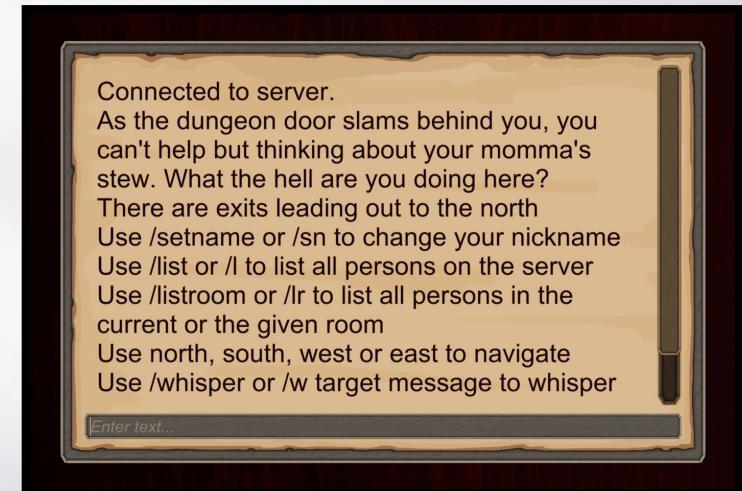
Recap 2 – What have we discussed so far?

- Network basics
- Application to application network communication over the internet



Today

- UDP & TCP characteristics in practice
- Some TCP/networking challenges
- Assignment 2 – Chatbox introduction



UDP vs TCP

UDP vs TCP demonstration

- The difference in *setup* between UDP & TCP

Difference in UDP/TCP setup?

- Checkout:
 - 001_basic_tcp_echo_server_client_commented (lecture 1) with
 - 001_basic_udp_echo_server_client_commented (lecture 2)
- Main thing to note:
 - Connection vs connectionless
 - 'Reading into a byte array' vs 'returning a byte array'
 - Ability to serve multiple clients, packets can come from anywhere
 - Duplicated packets? Missing packets? Unordered packets?

UDP vs TCP demonstration

- ~~The difference in *setup* between UDP & TCP~~
- The differences in *characteristics* between UDP & TCP
 - Packet loss
 - Packet duplication
 - Packet ordering

Difference in UDP/TCP characteristics?

- Hard to test with this setup:
 - Sending/receiving messages requires human intervention (type message, send, wait, repeat)
 - Perfect network conditions (everything runs on localhost ☺)
- Two changes:
 - a different test setup for auto send/receive
 - use a tool to simulate not so perfect network conditions

Test setup

- Again two examples:
 - 002_tcp_characteristics_demo
 - 002_udp_characteristics_demo
- Main idea: send increasing numbers and check order of arrival:
 - CS:1 → SR:1 → OK*
 - CS:2 → SR:2 → OK
 - CS:3 → SR:1 → Erm houston?
 - CS:4 → SR:4 → OMG, we've lost kenny!
 - CS:5 → SR:3 → Nevermind, we've found him!

(* CS = Client Send, SR = Server Receive)

Some details on the BitConverter class (MSDN)

- “The BitConverter class helps manipulate value types in their fundamental form, as a series of bytes. A byte is defined as an 8-bit unsigned integer. The BitConverter class includes static methods to convert each of the primitive types to and from an array of bytes” →
- You can play around with this [here](#)

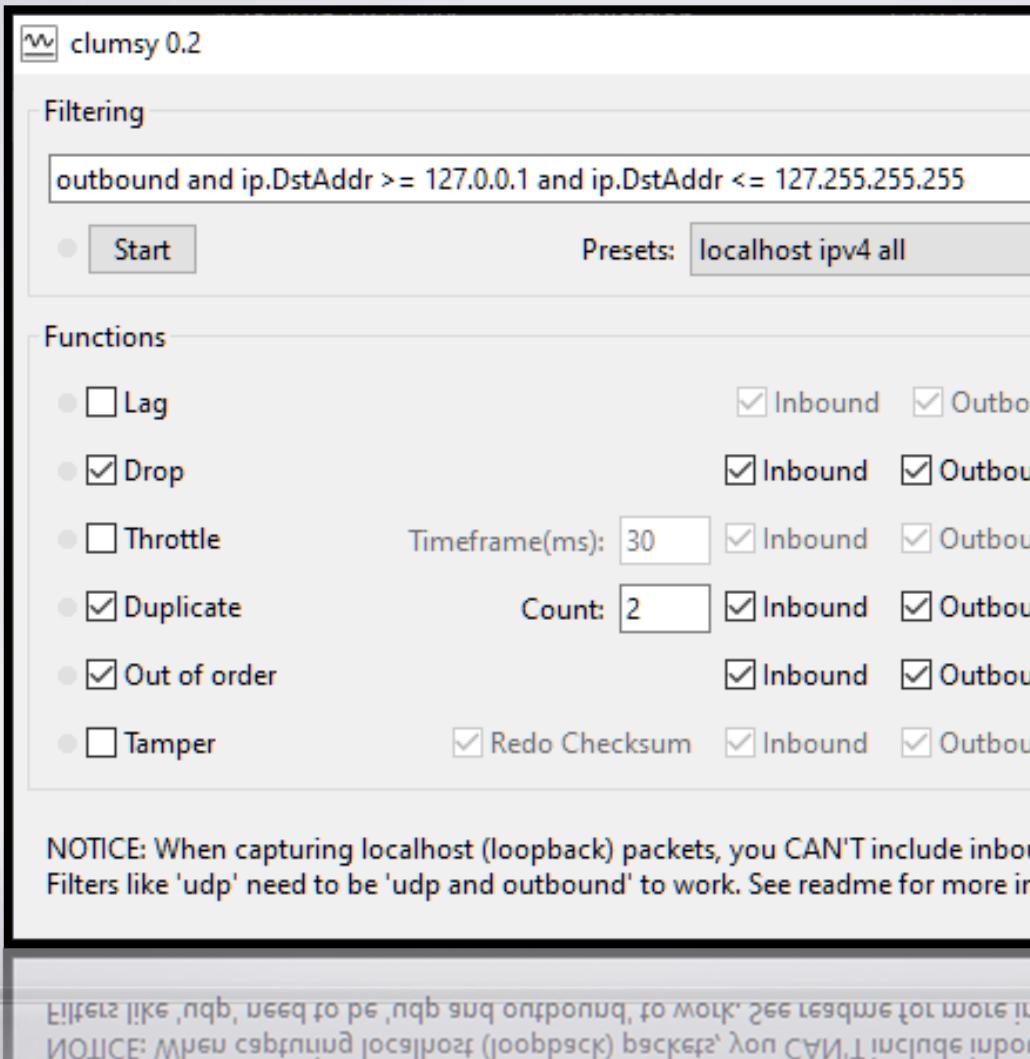
Type	To byte conversion	From byte conversion
Boolean	GetBytes(Boolean)	ToBoolean
Char	GetBytes(Char)	ToChar
Double	GetBytes(Double) -or- DoubleToInt64Bits(Double)	ToDouble -or- Int64Bits.ToDouble
Int16	GetBytes(Int16)	ToInt16
Int32	GetBytes(Int32)	ToInt32
Int64	GetBytes(Int64)	ToInt64
Single	GetBytes(Single)	ToSingle
UInt16	GetBytes(UInt16)	ToUInt16
UInt32	GetBytes(UInt32)	ToUInt32
UInt64	GetBytes(UInt64)	ToUInt64

Float	GetBytes(Float)	ToFloat
String	GetBytes(String)	ToHexString

Why do you need it?

- What is the size of the following objects in bytes?
 - `Int32.MaxValue` ($= 2,147,483,647 = 0x7F FF FF FF$) passed as a string?
 - `Int32.MaxValue` passed as an integer?
 - 255 passed as a string?
 - 255 passed as an integer?
- Conclusion:
 - integer byte size is fixed at 4 bytes → easier to handle
 - no string <-> integer conversion required → faster

Simulating bad network conditions with [Clumsy](#)



Word of warning:

Tools like this ***can*** break your system.

If you use it, make sure you:

- Use exactly the same settings (filter & functions)
- Turn the filtering **off** before closing the application or shutting down your pc.

99.9% the time nothing bad will happen, but I had to reinstall once 😊

Proves what we already knew...

UDP

- Connection-**less**
- Unreliable,
packets **can** be:
 - lost,
 - duplicated,
 - reordered

TCP

- Connection **oriented**
- Reliable,
packets **cannot** be:
 - lost,
 - duplicated,
 - reordered

TCP message boundaries

Playing around with the echo client/server



- Try this at home for both the UDP & TCP example:
 - Duplicate the line on the server that sends a reply
 - TCPEchoServerSample line 66:
`stream.Write(inBytes, 0, inByteCount);`
 - UDPEchoServerSample line 34:
`client.Send (inBytes, inBytes.Length, endPoint);`
 - Start sending messages from the client and see what happens

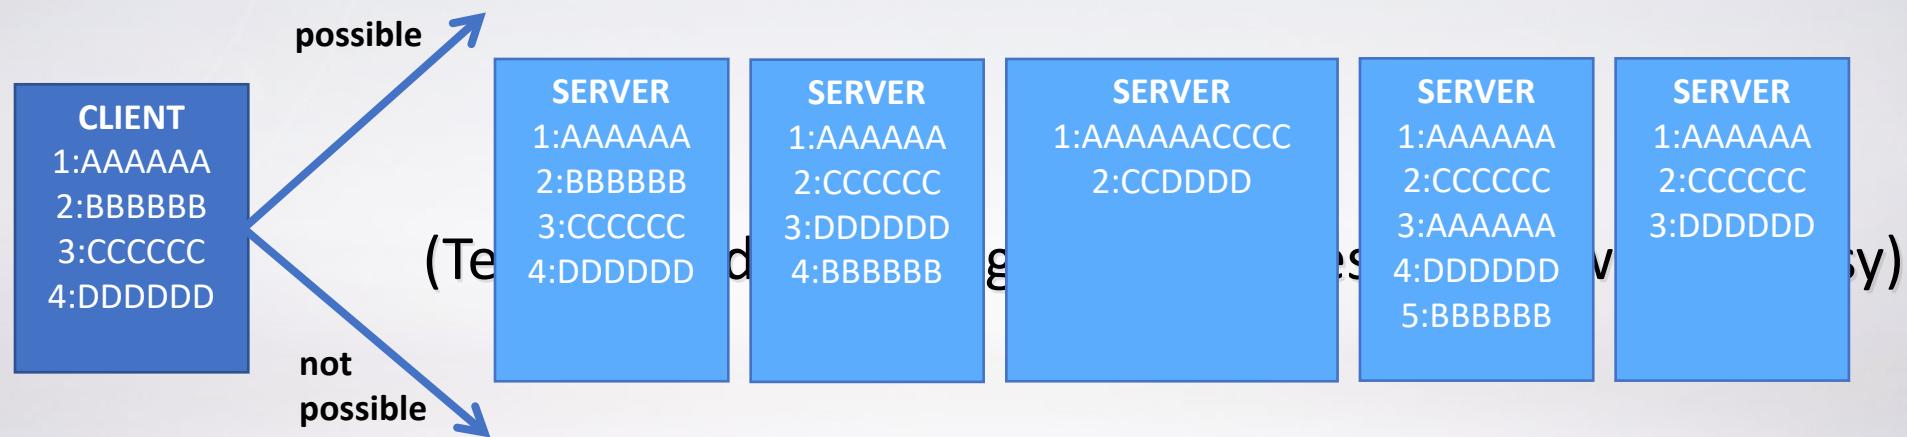
Conclusion

- TCP combines messages: $A\ B\ C\ D\ E \rightarrow A\ B\ C\ D\ E$
- UDP does *not* combine messages: $A\ B\ C\ D\ E \rightarrow A\ C\ E\ D$
- Another way to say this is that:
 - UDP *preserves message boundaries*
 - TCP does *not preserve message boundaries*

UDP preserves message boundaries

- UDP Packet sending and receiving is **all or nothing**
- *If* we send a packet of x bytes:
 - You might get the packet
 - You might not get the packet
 - You might get it twice
 - You might get packets out of order
 - But **if** it arrives.... it will be **exactly** x bytes...

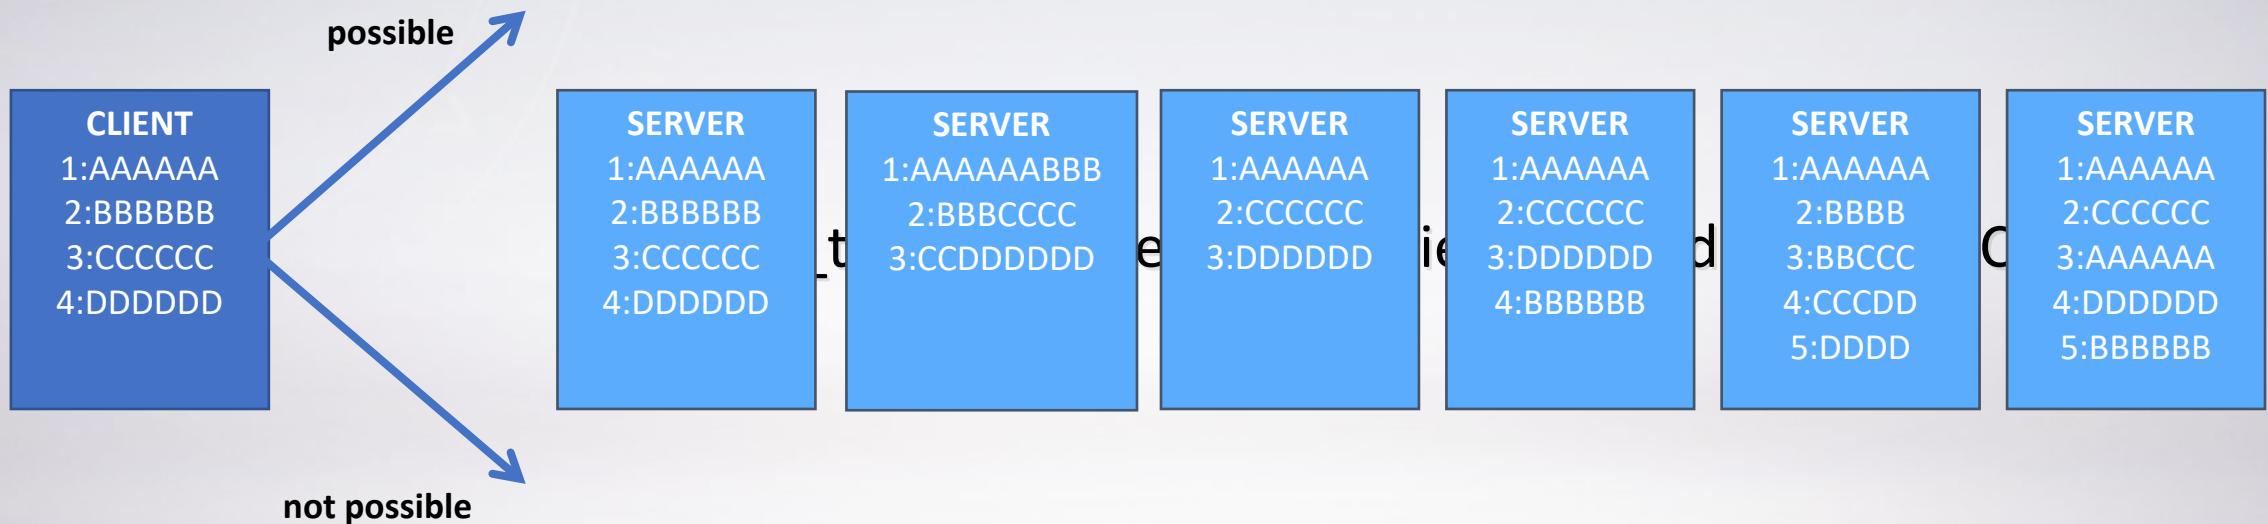
UDP preserves message boundaries



TCP does *not* preserve message boundaries

- As long as there is a connection TCP packets (bytes) ...
 - will arrive
 - in order
 - non duplicated
- But one send does not *necessarily* equal one receive:
 - One send *might* be one receive,
 - One send might be *multiples* receives,
 - Multiple sends might be *one* receive

TCP does not preserve message boundaries



Why should we care? Why are boundaries important?

Imagine a simple text protocol to move an object:
move,objectId,x,y e.g. move,playerOne,100,200

Pseudo parse code:

```
string[] command = message.split (",");  
if (command[0] == "move") {  
    FindObject (command[1]).SetXY (  
        int.parse (command[2]), int.parse (command[3])  
    );  
}
```

Life without boundaries

- Send 1 *move,playerOne,100,200*
 - Send 2 *move,playerTwo,150,300*
-

- Receive 1 *move,playerOne,10*
- Receive 2 *0,200move,playerTwo,150,*
- Receive 3 *300*

- Good luck parsing this crap!
- Imagine if these were byte arrays representing objects (aka lecture 3)

Why does TCP work like this?

- TCP packets are ordered, non duplicate & reliable
- The only way to implement that is by keeping track of what you have:
 - sent – so that you can resend data, if necessary
 - received – so that you can order & filter new data
- Done through underlying send & receive buffers
 - One side writes data into the buffer
 - Other side reads data from the buffer
 - OS's on both sides send/receive when they 'feel' like it
 - These two buffers might or might not be in sync
- So again: one send does not *necessarily* equal one receive

Inner workings of Stream.Write ?

- void Stream.Write (buffer, offset, size):
 - See if there is room in the send buffer, if so write, if not block (!)

Inner workings of Stream.Read ?

- int Stream.Read (buffer, offset, size):
 - Checks if there is data in the Read buffer
 - might return:
 - size (exactly the amount of bytes you asked for, if present)
 - ***less than*** size (even down to **only returning 1 byte** at a time!)
 - 0 (0 means: 'reached end of connection' and **not** 'no data available')
 - an error
 - might block *until* there is data available
- In other words:
 - Read blocks until some **but not necessarily all** data is available
 - You cannot just call Read (... , ..., x) and expect it to return x bytes.

Conclusion

UDP

- Connection-**less**
- Unreliable,
packets **can** be:
 - lost,
 - duplicated,
 - reordered
- Preserves message boundaries
- (Max 64k message size)

TCP

- Connection **oriented**
- Reliable,
packets **cannot** be:
 - lost,
 - duplicated,
 - reordered
- Doesn't preserve message boundaries
- ("Unlimited" message size)

TCP Message Boundary Detection

Boundary implementation options

- Language independent “options”
 - Use fixed size messages
 - Delimit each message with a token
 - Sending the message size before actual message (*this is the best option!*)
- C# specific options:
 - BinaryFormatter, BinaryReader/Writer, StreamReader/Writer
(but these also send/receive the message size below the surface)

Best option: Send message size before actual message!

- Principle:
 - Sender: *Send* msg size first, then message itself
 - Receiver: *Read* msg size first, then *loop Read until whole message is received*
- For example, to send/receive "hello":
 - Sender side:
 - Send message length → Send int 5 (4 bytes)
 - Send message itself → Send "hello" (5 bytes)
 - Receiver side:
 - Read message length (4 bytes) → the number 5
 - Read 5 bytes and convert to string → "hello"

Requirements

1. A way to send the message size as a **fixed amount** of bytes (why??).
2. A way to correctly read exactly x bytes.

Requirement 1 Sending size as a fixed amount of bytes

- If the message size is described by a variable amount of bytes, we still don't know how many bytes to read!
- This means we cannot use strings to send the message size: "1", "12", "154" → all different amount of bytes
- Solution:
 - use BitConverter class instead and convert int's to bytes instead of strings

Requirement 2 Reading x bytes correctly

1.500.000 (составляется == председатель) : председатель

Adding some utility methods...

```
public static void SendMessage (NetworkStream pStream, byte[] pMessage) {
    pStream.Write (BitConverter.GetBytes (pMessage.Length), 0, 4);
    pStream.Write (pMessage, 0, pMessage.Length);
}

public static byte[] ReceiveMessage (NetworkStream pStream) {
    int byteCountToRead = BitConverter.ToInt32(ReadBytes (pStream, 4), 0);
    return ReadBytes (pStream, byteCountToRead);
}

public static void SendString (NetworkStream pStream, string pMessage, Encoding pEncoding) {
    SendMessage (pStream, pEncoding.GetBytes (pMessage));
}

public static string ReceiveString (NetworkStream pStream, Encoding pEncoding) {
    return pEncoding.GetString (ReceiveMessage(pStream));
}
```

}
(Test 004_tcp_message_boundaries_correct_demo with Clumsy)

Conclusion

- TCP provides a reliable mechanism to send/receive unique messages in order, and with some helper classes (custom or provided) we can “fix” the fact that it doesn’t preserve message boundaries by default.
- Examples of existing helper .Net classes which you could also use:
 - BinaryReader/Writer
 - StreamReader/Writer
 - Serializers such BinaryFormatter/Xml/Json

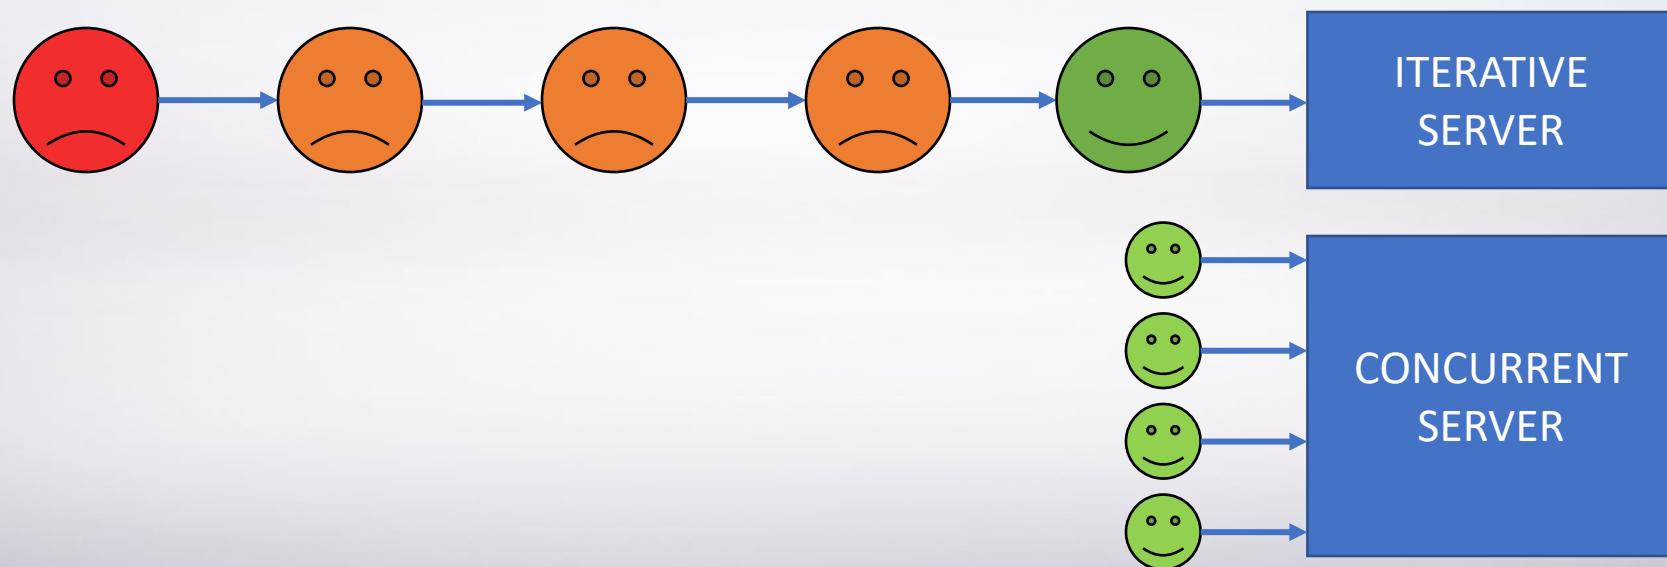
Blocking calls vs multiple clients

Serving multiple clients

- All examples showed 1 client/1 server
- Client 1 kept looping & the server kept serving...
- What will happen when client 2 arrives?
 - Client 2 will be put in a queue until Client 1 is done
- When is client 1 done?
 - When either the client or the server closes the connection!

Iterative vs Concurrent server

- A server that puts client 2 on hold while serving client 1 is called an **iterative** server
- A server that is able to serve multiple clients at the “same” time is called a **concurrent** server



How to move to a 'concurrent' server setup?

- Possible solutions:
 - Threading (manual thread creation, asynchronous callback methods etc)
 - Using sockets in non blocking mode, Select(...)
 - Timeouts & error handling
 - Polling
- Polling is by far the easiest solution: check if you **would** block **before** actually blocking, using:
 - if (listener.Pending())
 - if (client.Available > 0)

Example 005_concurrent_tcp_echo_server

- Things to note:
 - Both client and server are using StreamUtil
 - Server serves multiple clients at 'once' by polling before blocking
 - Both client & server are still single threaded
 - No communication between clients
 - Assumes if TCPClient.Available > 0 a whole message is available.
If a client only sends a message size and then quits without sending the rest, everything still blocks.
 - Disconnected clients are not detected (yet)
 - No error handling

Disconnecting clients & error handling

Error handling

- None of the examples showed error handling, while most networking calls can throw exceptions
- A crashing client is not so bad, but a server that crashes because the client crashed is pretty bad...
- How can we protect against that?
 - Error handling, e.g. try-catch (of course)
 - Remove faulty clients (otherwise errors keep happening)
 - Exceptions might keep occurring
 - We might keep waiting for a client to act indefinitely

Detecting faulty (disconnected etc) clients...

- Three different situations:
 - Client calls Close on Socket/Stream
 - Client shuts down without calling Close
 - Physical channel goes down (eg you pull the network plug)
- How can we detect this?
 - There is a so called TcpClient.Connected flag,
but that only updates **after** sending
 - In other words: the flag we would like to use to prevent problems,
only updates after a problem has already occurred

Handling errors & disconnects

- Good enough for now:
 - Try to send/receive data
 - Try catch any errors that occur (and print them!)
 - Remove clients whose Connected property is false
- In practice:
 - connections issues might take a long time to detect
 - this means using a heartbeat system would be even better
 - send out heartbeat every x seconds
 - keep track of last heartbeat received from other side
 - no heartbeat received for y seconds from other side? Assume painful death, exit.

"Evil" clients

What are the simple things evil clients could do?

- Clients could overload the server with messages, forcing the server to spin forever trying to process them
- Clients could ‘lie’ about their messages:
 - “I am going to sent you 5 bytes!”
 - *“Please hold while I am secretly not sending anything ghe ghe ghe”*
- Clients can stop reading from the buffer forcing server to block on send
- Clients could tamper with your messages (Clumsy / decompiling / etc)
- Although these are all 'fixable', assume ‘nice’ clients for now:
 - Clients might join & leave (and you are not allowed to crash due to that)
 - Clients will not intentionally try to stop/crash your server

Summing up ...

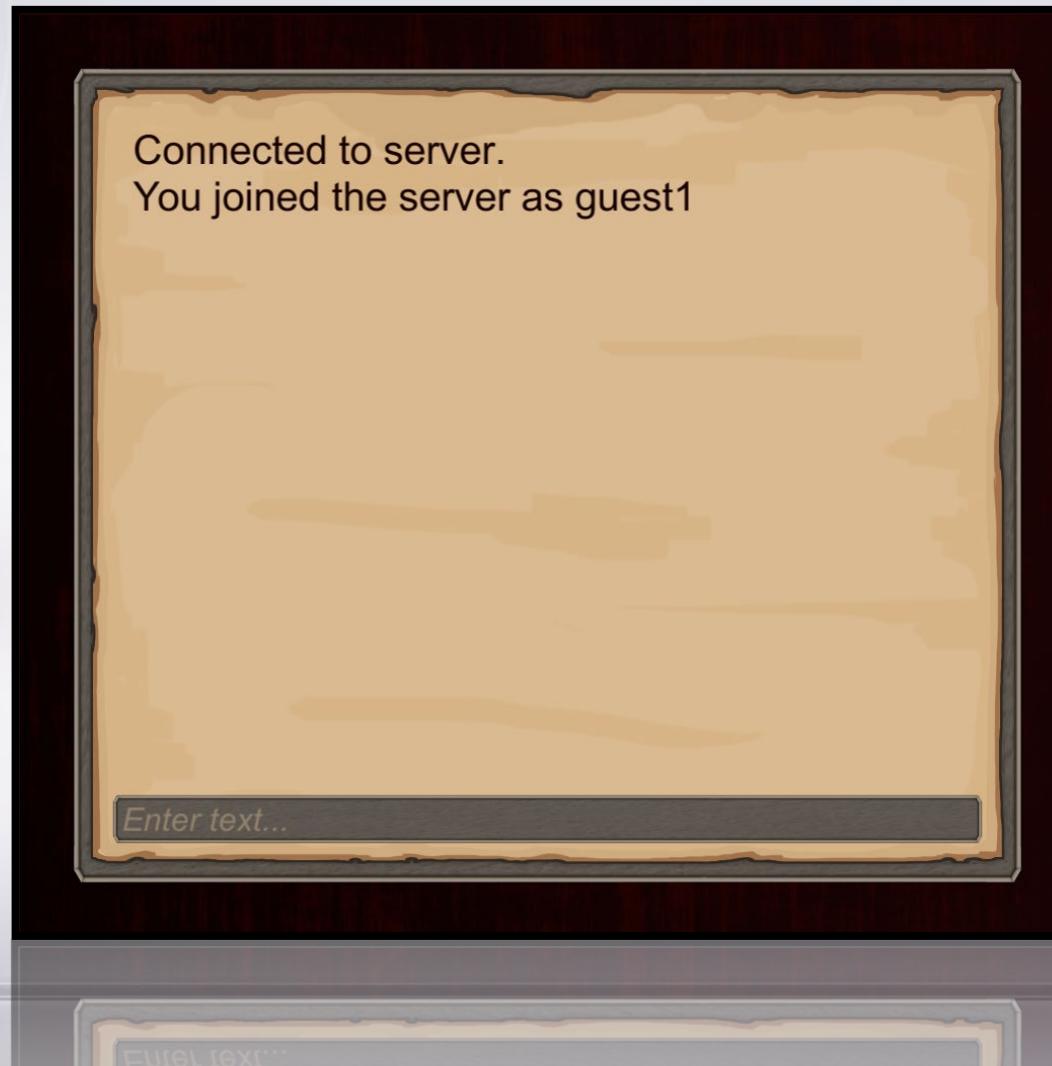
What have we seen today?

- The difference between UDP & TCP in practice and how to test your code under bad networking conditions (Example 001 & 002)
- How UDP & TCP deal with message boundaries and how we can deal with them ourselves (Example 003 & 004)
- How we can implement a concurrent server by polling (Example 005)
- Some general tips on dealing with (faulty, crashing & evil) clients
- In other words... everything you need for...

Introducing assignment 2

Implementing a chatbox in Unity

Assignment 2 – Chatbox



Assignment 2 – Implementing a chatbox

- Requirements plus starting code on BlackBoard
- Main difference with command line Tcp echo client server demo?
 - Client & server can no longer block
 - Incoming messages should be bounced to all connected users
 - Users need to be identified through their TcpClient handle
 - Messages can have meaning (setname, whisper, etc)
- Check **both** grading form & assignment requirements!
 - For example: Server is not allowed to crash (not explicitly mentioned in assignment)
- 99% similar to assignment 1 last year:
 - 'Getting started' section added
 - Minor details changed: be sure to double check if you are doing the redo!

Non reproducible errors are hard to fix

- Programs that are non deterministic
- Programs that run in different threads
- Programs that depend on multiple cooperating processes/external data
- Programs that depend on best-effort protocols
- In other words:
 - (threaded) client/server programs can be a nightmare to debug 😊
 - complete your implementation step by baby step
 - print logical debug info wherever possible

